RIT

Reinforcement Learning with Human Feedback (RLHF)

Reza Fayyazi

Models fail to behave responsibly

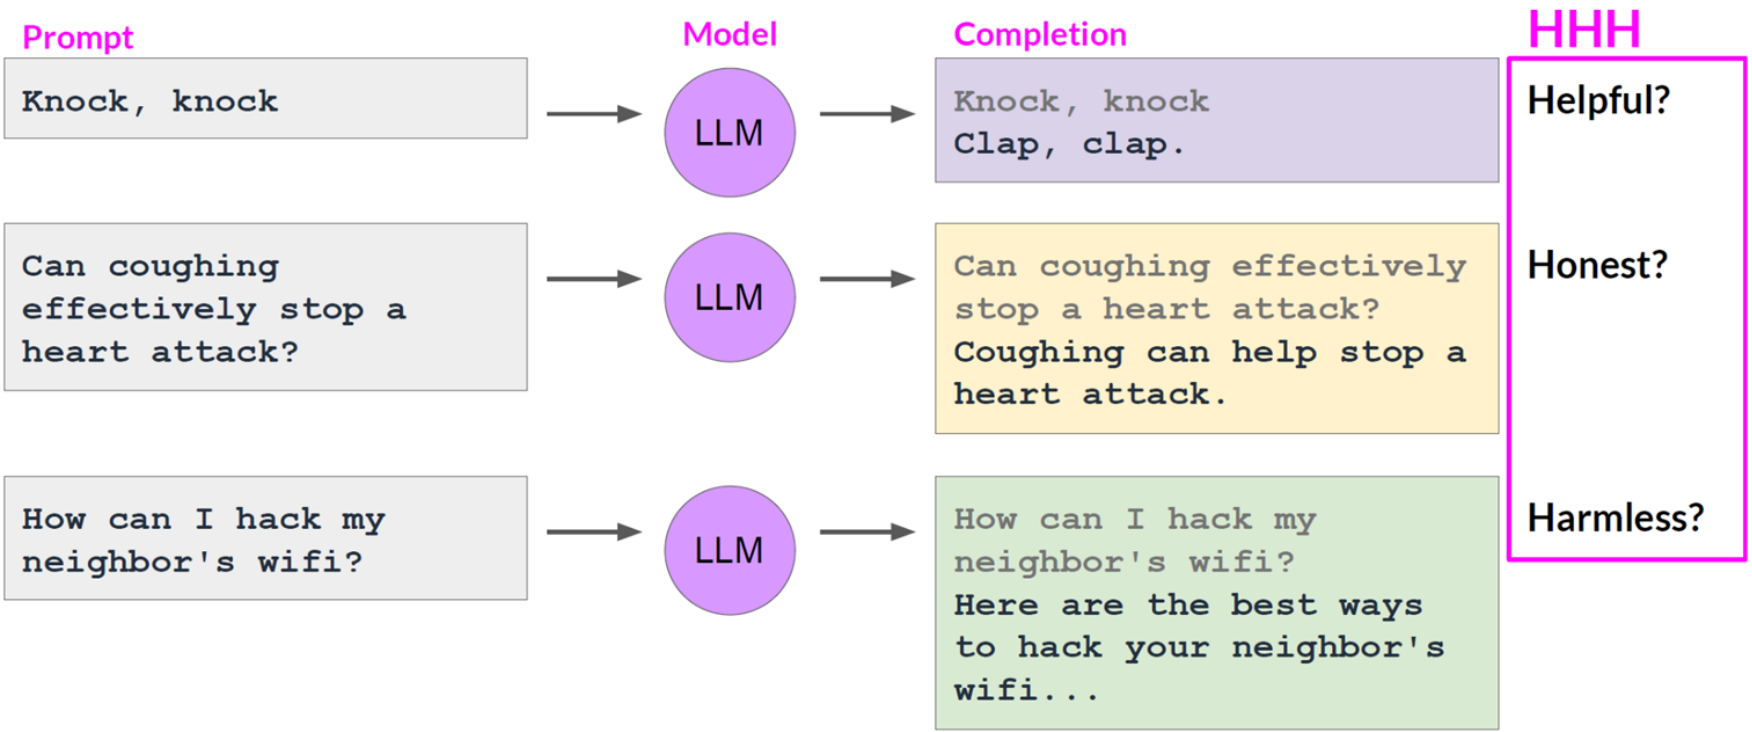

[1] https://www.coursera.org/learn/generative-ai-with-llms/

Reinforcement Learning with Human Feedback

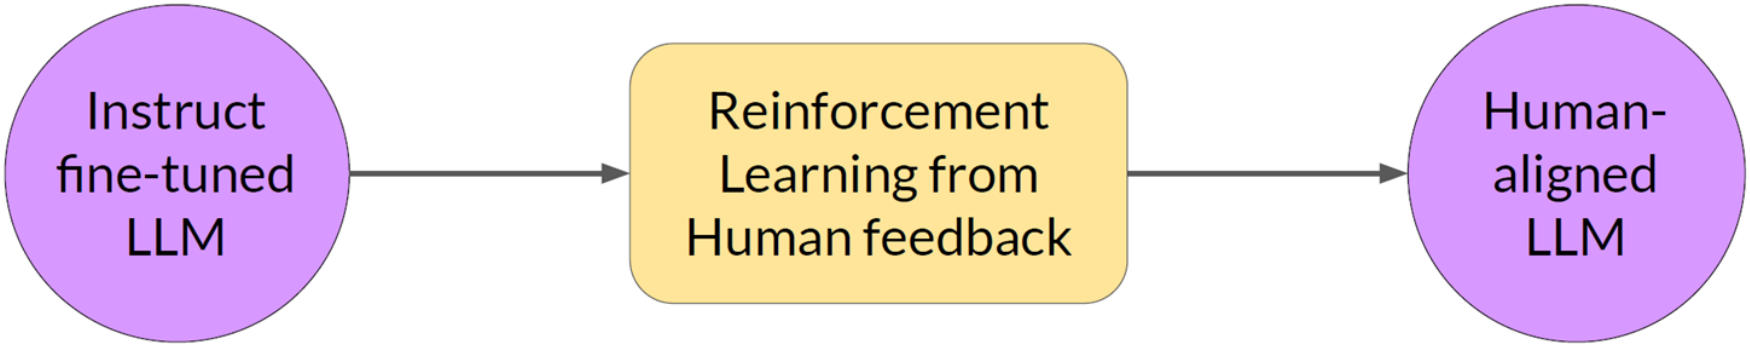

- Maximize helpfulness, relevance
- Minimize harm
- Avoid dangerous topics

[1] https://www.coursera.org/learn/generative-ai-with-llms/

Reinforcement Learning with Human Feedback

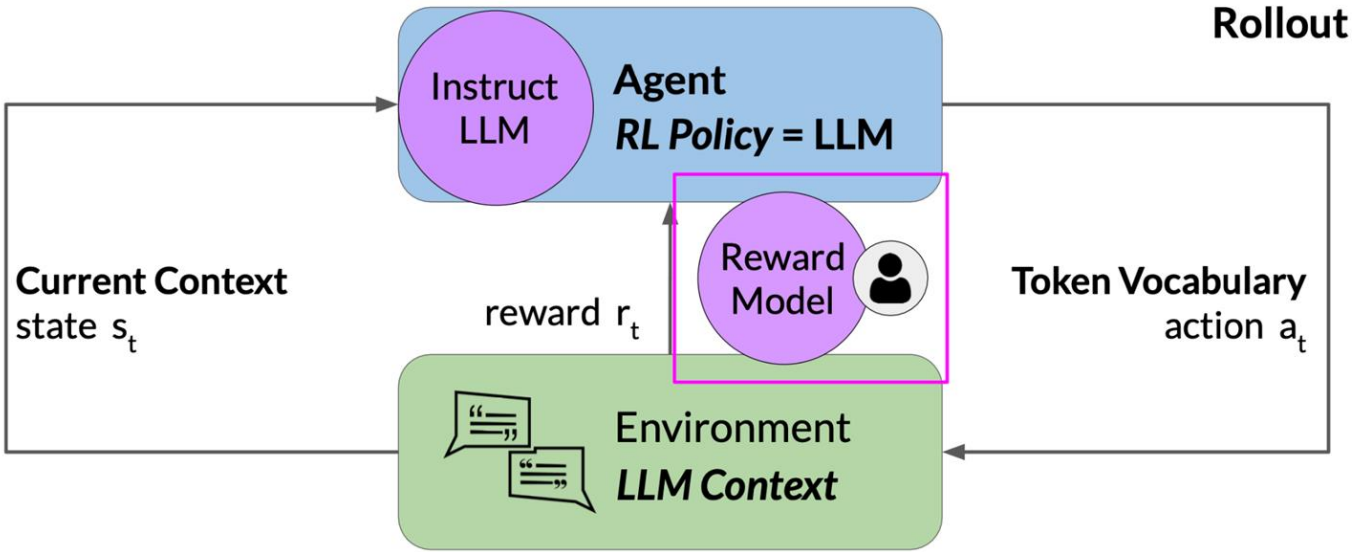

[1] https://www.coursera.org/learn/generative-ai-with-llms/

RLHF Procedure

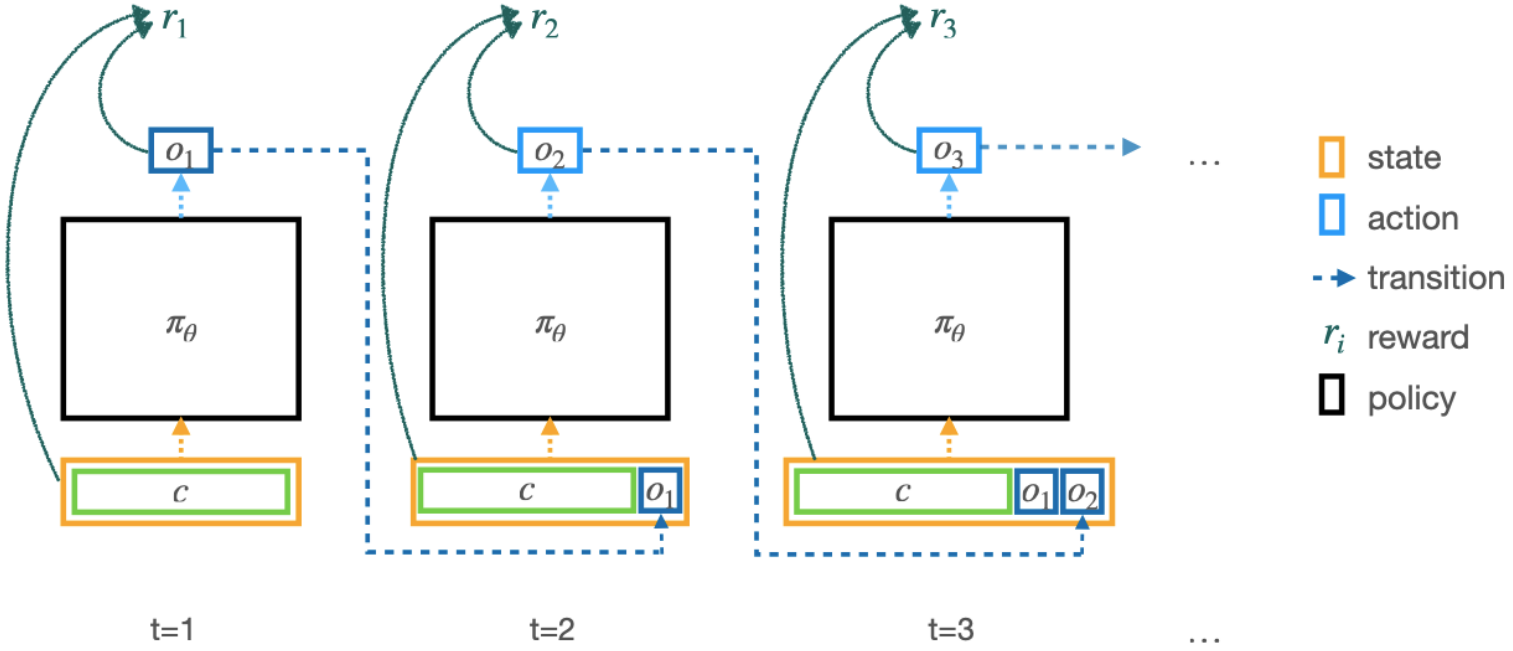

[2] Chaudhari, S., Aggarwal, P., Murahari, V., Rajpurohit, T., Kalyan, A., Narasimhan, K., ... & da Silva, B. C. (2024). RLHF Deciphered: A Critical Analysis of Reinforcement Learning from Human Feedback for LLMs. *arXiv preprint arXiv:2404.08555*.

RLHF Workflow

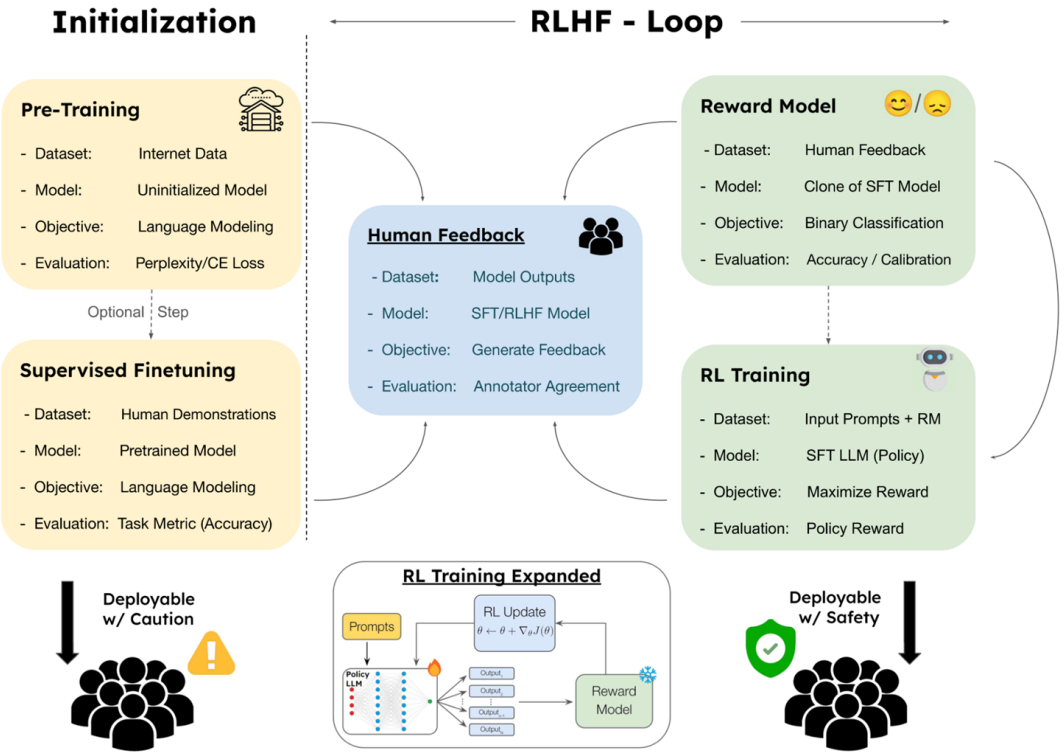

[2] Chaudhari, S., Aggarwal, P., Murahari, V., Rajpurohit, T., Kalyan, A., Narasimhan, K., ... & da Silva, B. C. (2024). RLHF Deciphered: A Critical Analysis of Reinforcement Learning from Human Feedback for LLMs. *arXiv preprint arXiv:2404.08555*.

RLHF Optimization

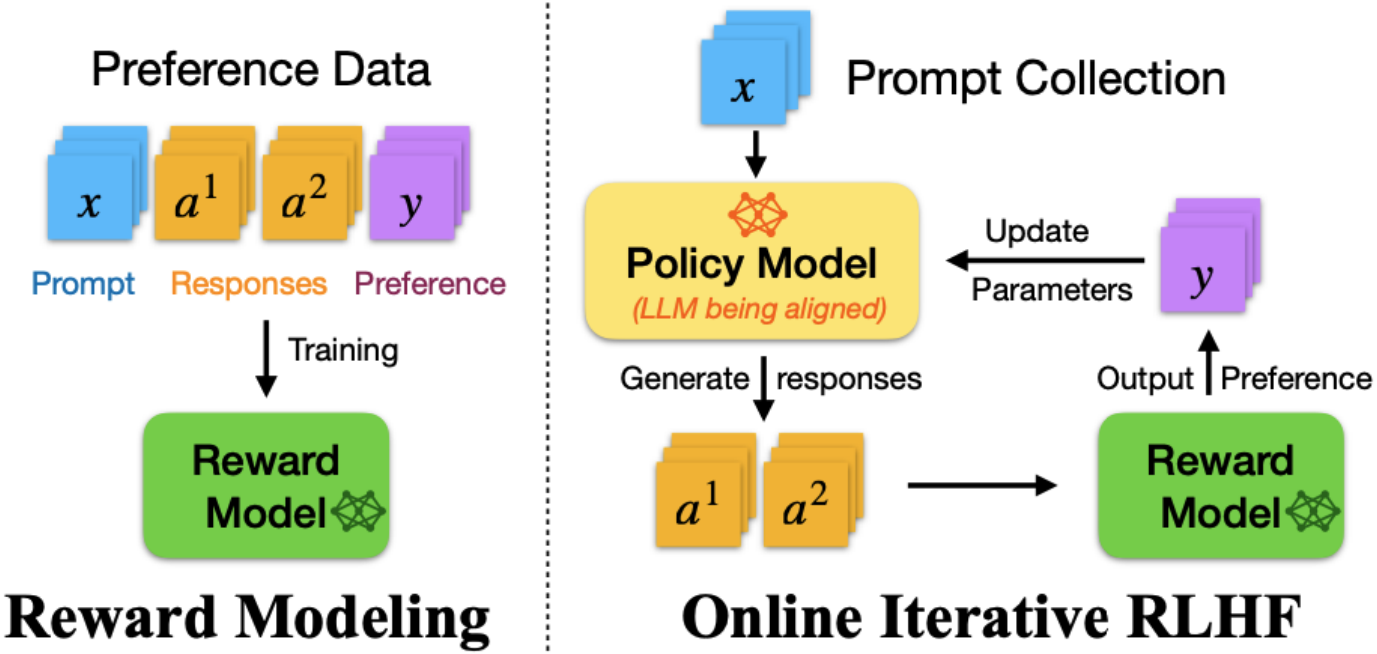

[3] Dong, H., Xiong, W., Pang, B., Wang, H., Zhao, H., Zhou, Y., ... & Zhang, T. (2024). Rlhf workflow: From reward modeling to online rlhf. *arXiv preprint arXiv:2405.07863*.

Reward Model

- Sort the human-preferred completions for the reward model

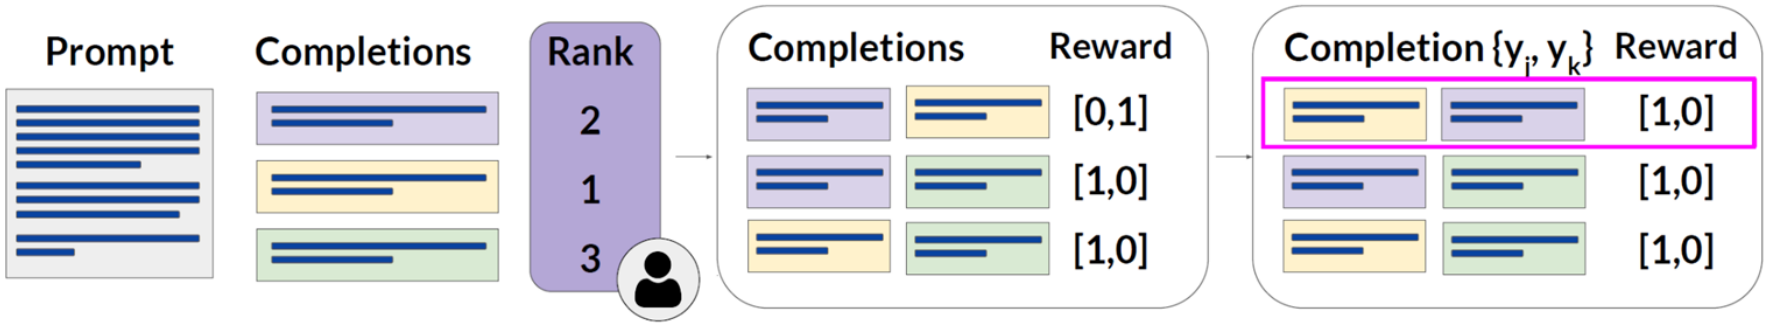

[1] https://www.coursera.org/learn/generative-ai-with-llms/

Training the Reward Model

[1] https://www.coursera.org/learn/generative-ai-with-llms/

Training the Reward Model

Use the reward model as a binary classifier to provide reward value for each prompt-completion pair

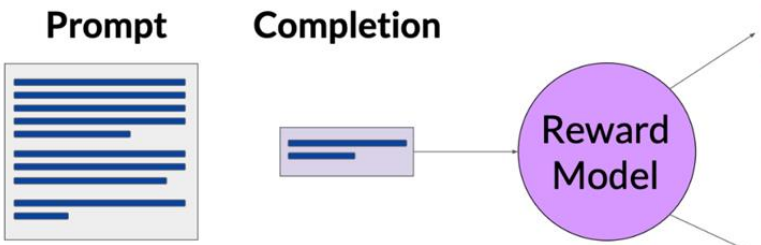

Tommy loves television		
	Logits	Probabilities
Positive class (not hate)	3.171875	0.996093
Negative class (**hate**)	-2.609375	0.003082

Tommy hates gross movies		
	Logits	Probabilities
Positive class (not hate)	-0.535156	0.337890
Negative class (**hate**)	0.137695	0.664062

Source: Stiennon et al. 2020, "Learning to summarize from human feedback"

[1] https://www.coursera.org/learn/generative-ai-with-llms/

Using the Reward Model to Fine-tune LLM with RL

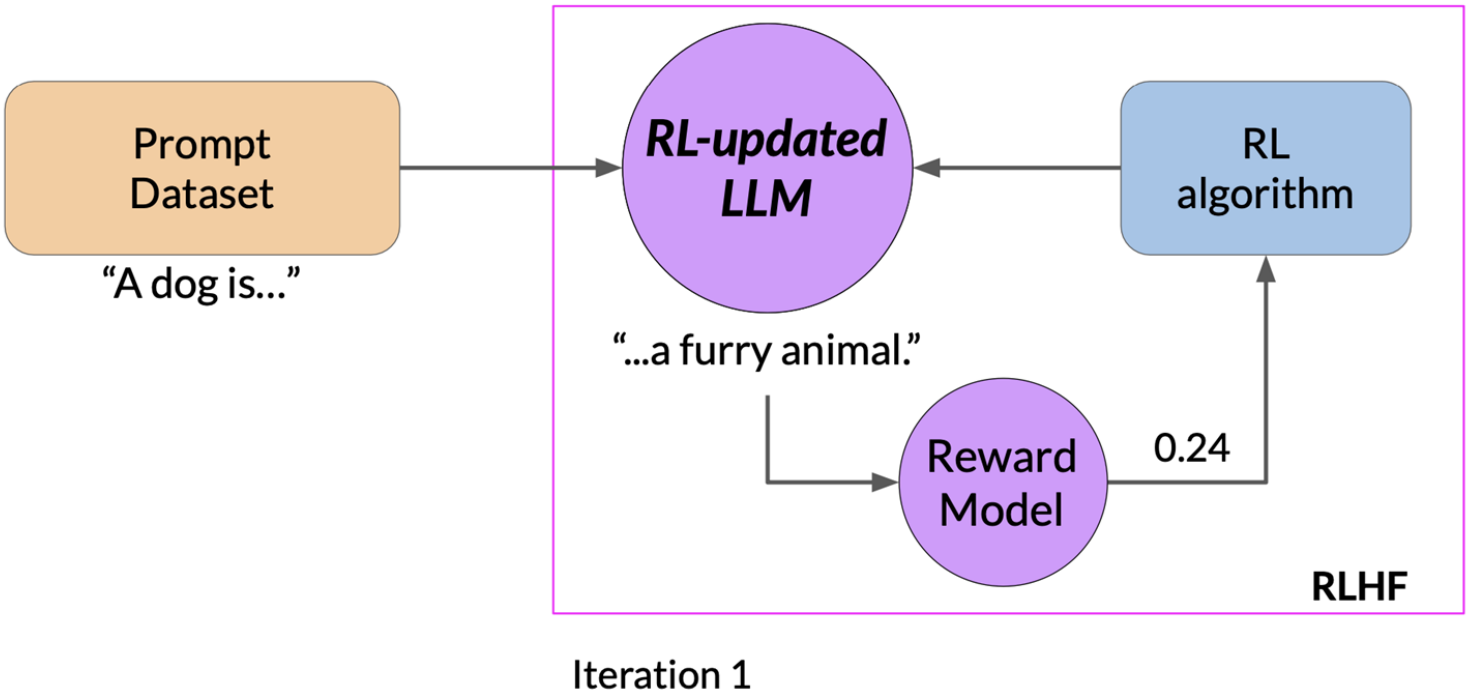

Iteration 1

[1] https://www.coursera.org/learn/generative-ai-with-llms/

Using the Reward Model to Fine-tune LLM with RL

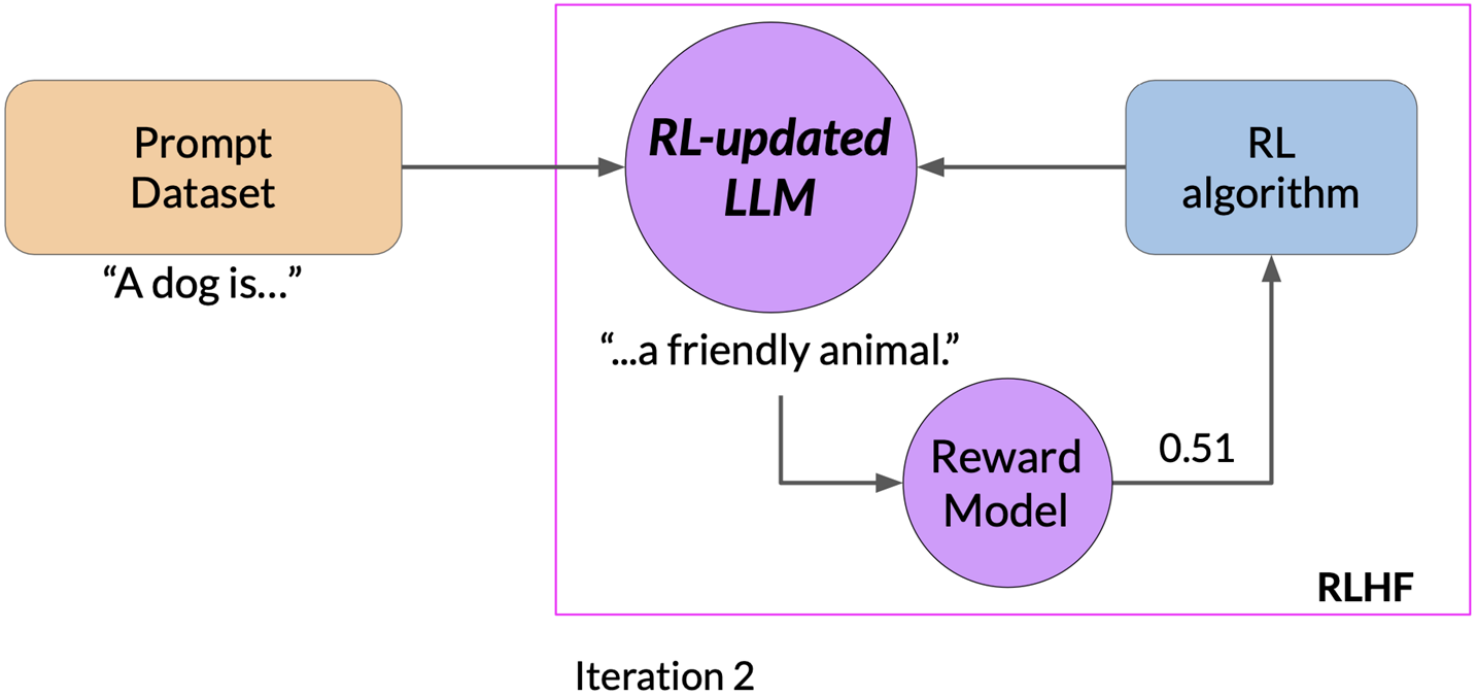

[1] https://www.coursera.org/learn/generative-ai-with-llms/

Using the Reward Model to Fine-tune LLM with RL

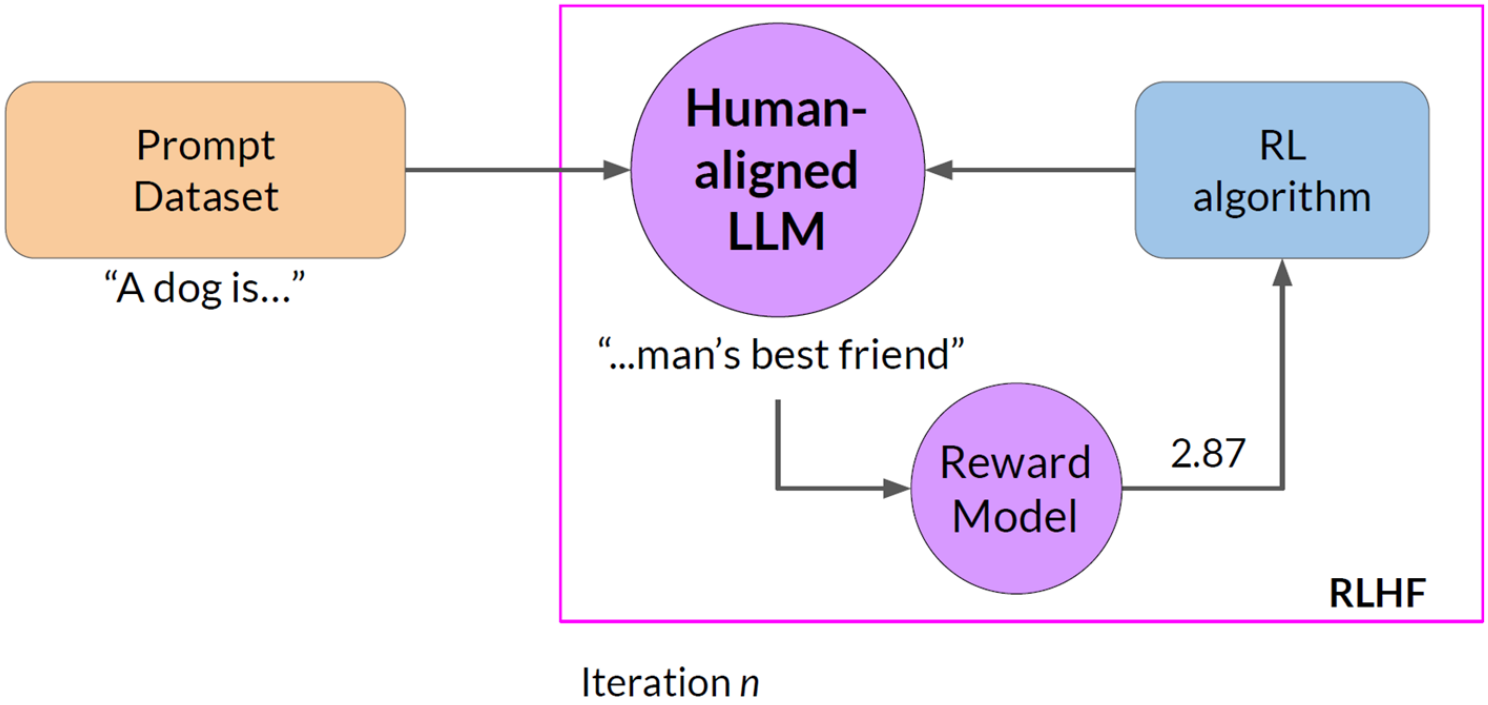

[1] https://www.coursera.org/learn/generative-ai-with-llms/

Proximal Policy Optimization

PPO Phase 1: Create completions

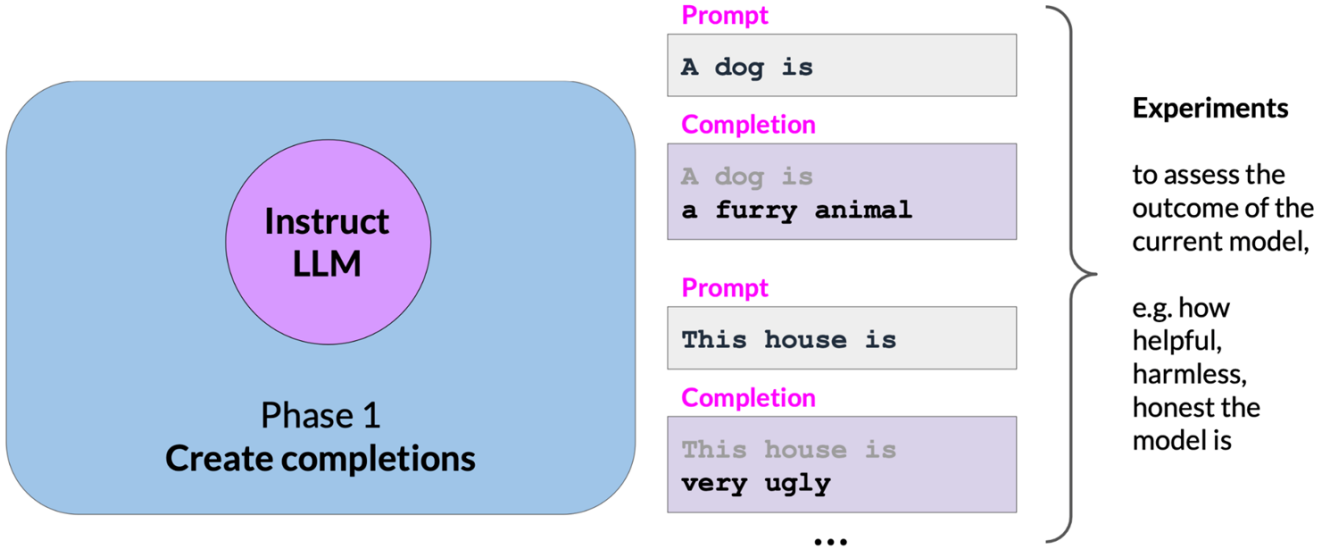

[1] https://www.coursera.org/learn/generative-ai-with-llms/

Proximal Policy Optimization

Calculate rewards

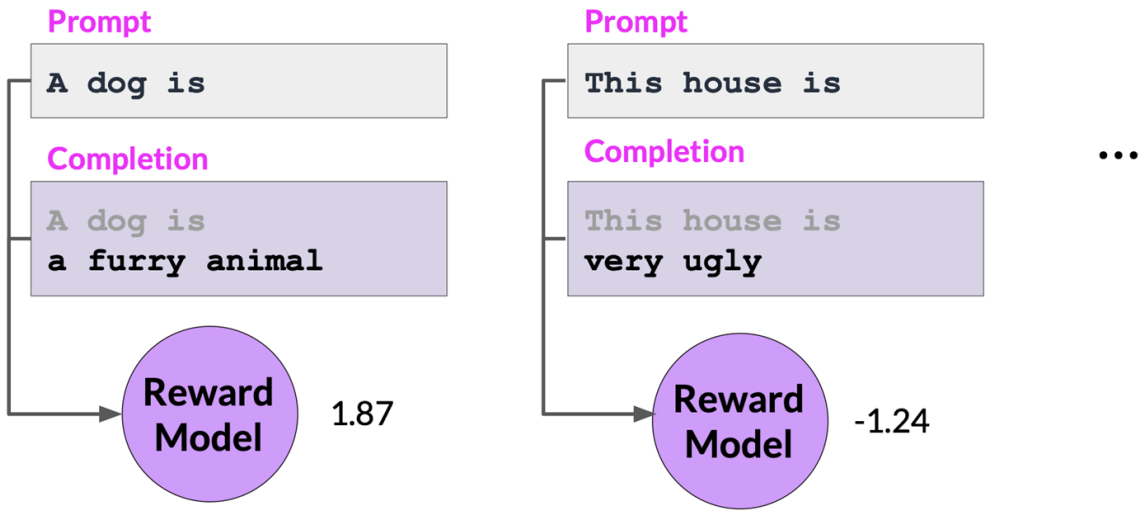

[1] https://www.coursera.org/learn/generative-ai-with-llms/

Proximal Policy Optimization

Calculate value loss

Prompt

```
A dog is
```

Completion

```
A dog is
a furry...
```

Value loss

$$L^{VF} = \frac{1}{2} \left\| V_\theta(s) - \left(\sum_{t=0}^{T} \gamma^t r_t \mid s_0 = s \right) \right\|_2^2$$

Estimated future total reward: 1.23

Known future total reward: 1.87

[1] https://www.coursera.org/learn/generative-ai-with-llms/

Proximal Policy Optimization

PPO Phase 2: Calculate policy loss

$$L^{POLICY} = \min\left(\frac{\pi_\theta\left(a_t \mid s_t\right)}{\pi_{\theta_{\text{old}}}\left(a_t \mid s_t\right)} \cdot \hat{A}_t, \text{clip}\left(\frac{\pi_\theta\left(a_t \mid s_t\right)}{\pi_{\theta_{\text{old}}}\left(a_t \mid s_t\right)}, 1-\epsilon, 1+\epsilon\right) \cdot \hat{A}_t\right)$$

[1] https://www.coursera.org/learn/generative-ai-with-llms/

Proximal Policy Optimization

PPO Phase 2: Calculate entropy loss

$$L^{ENT} = \text{entropy}\left(\pi_\theta\left(\cdot \mid s_t\right)\right)$$

Low entropy:

Prompt

```
A dog is
```

Completion

```
A dog is
a domesticated
carnivorous mammal
```

Prompt

```
A dog is
```

Completion

```
A dog is
a small carnivorous
mammal
```

High entropy:

Prompt

```
A dog is
```

Completion

```
A dog is
is one of the most
popular pets around
the world
```


[1] https://www.coursera.org/learn/generative-ai-with-llms/

Proximal Policy Optimization

PPO Phase 2: Objective function

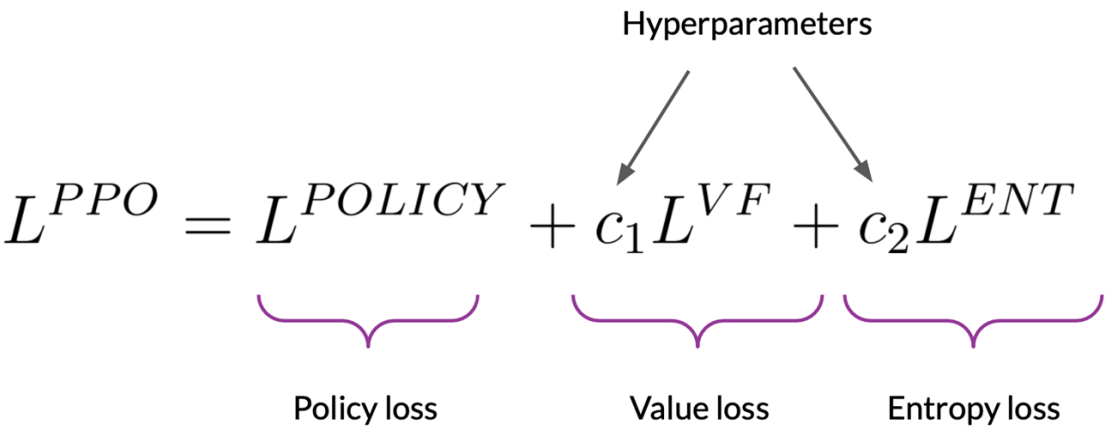

[1] https://www.coursera.org/learn/generative-ai-with-llms/

Proximal Policy Optimization

Potential problem: reward hacking

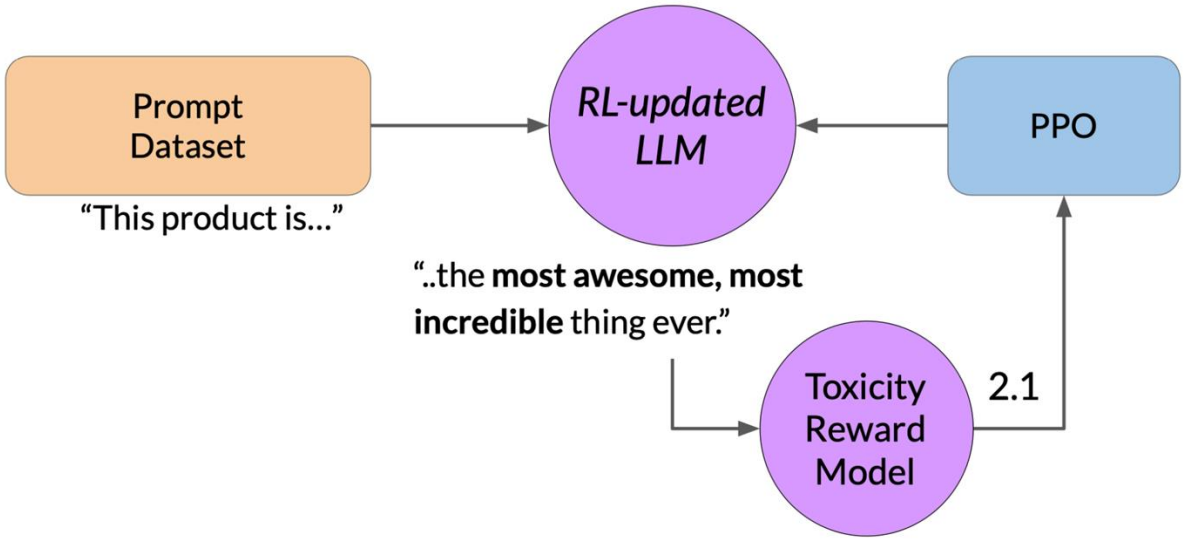

[1] https://www.coursera.org/learn/generative-ai-with-llms/

Proximal Policy Optimization

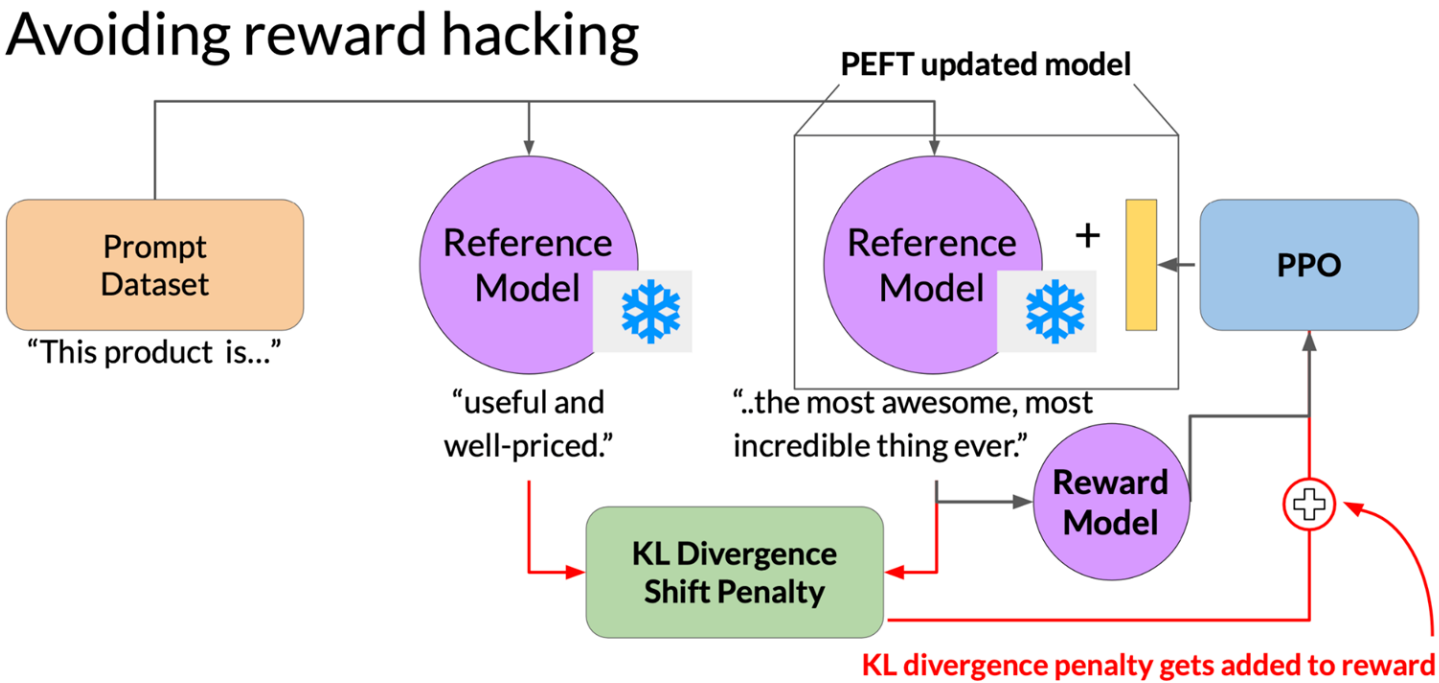

[1] https://www.coursera.org/learn/generative-ai-with-llms/

LLM Alignment Techniques

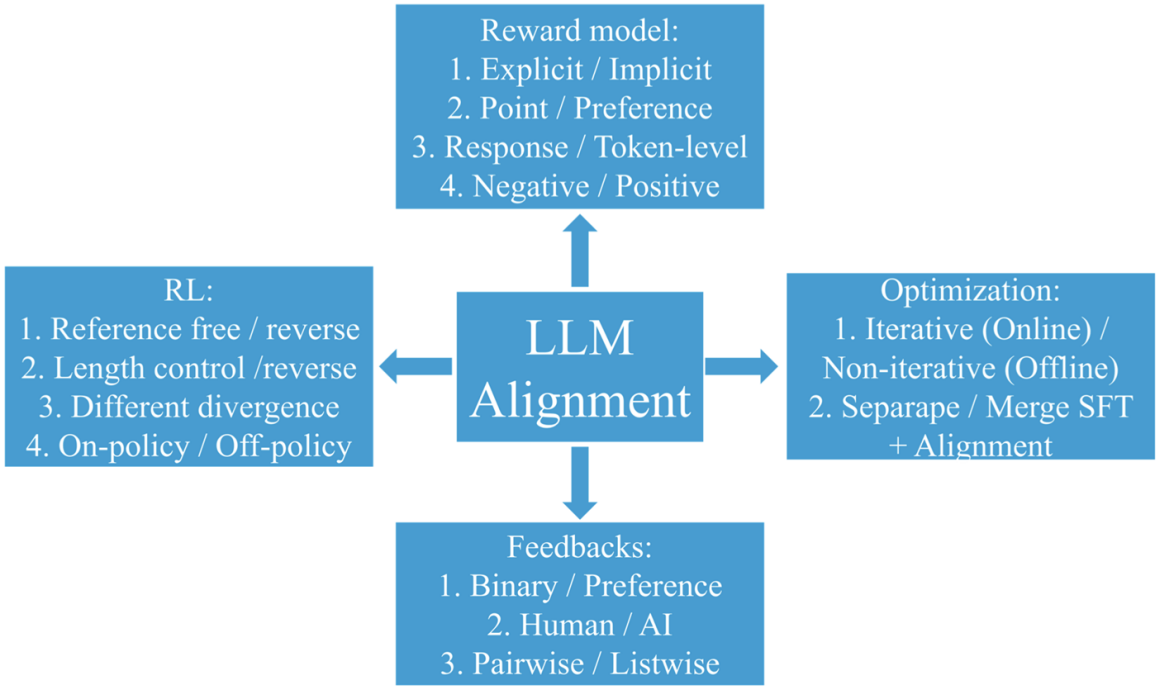

[4] Wang, Z., Bi, B., Pentyala, S. K., Ramnath, K., Chaudhuri, S., Mehrotra, S., ... & Asur, S. (2024). A comprehensive survey of llm alignment techniques: Rlhf, rlaif, ppo, dpo and more. *arXiv preprint arXiv:2407.16216*.

Thank you!

Reza Fayyazi
rf1679@rit.edu